

Natural language processing

Text as data: how machines read


POLI 158: AI & Politics · Session 5

Yiqing Xu


TODAY'S QUESTION

If an AI learned only from censored text, what would it believe, and could you tell?

The eight Wednesday sessions



Today in three parts



Where computers read text

Everyday life



autocomplete



translation



search



chatbots

Research



sorting posts by stance



placing parties left to right



finding what censors delete

Grimmer and Stewart (2013); King, Pan, and Roberts
(2013)

Politics and policy



screening hate speech




voting advice




sorting public comments

From text to numbers




Three concepts in today's video




Tokens

text becomes numbers



Embeddings


a map of meaning



Attention

context changes meaning

2 Embeddings




One sentence in five languages

Language	Universal Declaration of Human Rights, Article 1, first sentence	Tokens
English	All human beings are born free and equal in dignity and rights.	?
Chinese	人人生而自由，在尊严和权利上一律平等。	?
Hindi	सभी मनुष्यों को गौरव और अधिकारों के मामले में जन्मजात स्वतन्त्रता और समानता प्राप्त है ।	?
Burmese	လူတိုင်းသည် တူညီ လွတ်လပ်သော ဂုဏ်သိက္ခာဖြင့် လည်းကောင်း၊ တူညီလွတ်လပ်သော အခွင့်အရေးများဖြင့် လည်းကောင်း၊ မွေးဖွားလာသူများ ဖြစ်သည်။	?
Amharic	የሰው፡ፊጅ፡ሁሉ፡ሲወለድ፡ነፃና፡በክብርና፡በመብትም፡እኩልነት፡ያለው፡ነው።	?

👉 Guess which line needs the most tokens. 🖱️ Then try platform.openai.com/tokenizer.

Source: United Nations (1948), Universal Declaration of Human Rights. English and Chinese: official UN texts. Hindi, Burmese, Amharic: translations from the UN human rights office's translation project, as encoded by UDHR in Unicode. Petrov et al. (2023).

Stereotypes in word embeddings




 People: how quickly they pair two kinds of words (the Implicit Association Test).

 Embedding: how close the two kinds of words sit on the map.

Source: Caliskan, Bryson, and Narayanan (2017), *Science*, Table 1 (five of its ten tests). The authors do not claim the two sets of numbers are directly comparable: read the directions, not the gaps.

The transformer



► GPT: generative pre-trained transformer.

Source: Vaswani et al. (2017); Brown et al. (2020), Table 2.1; 3Blue1Brown (2024); the 2025 course slides.

Activity: tokens and word embeddings


Activity, part 1: tokens

In pairs, one laptop per pair. No login needed.

 platform.openai.com/tokenizer

1. Type a sentence in another language you know, then in English. Which needs more tokens?
2. Click the older tab, “GPT-4 & GPT-3.5 (legacy).” Does the gap grow or shrink?
3. Type your name and your hometown. Which break into pieces?
4. Type “strawberry.” Why might a chatbot miscount its r’s?

Screenshot: OpenAI’s tokenizer page, default tab, October 6, 2026.



The screenshot shows the OpenAI tokenizer interface. At the top, there are three tabs: "GPT-5.x & O1/3" (selected), "GPT-4 & GPT-3.5 (legacy)", and "GPT-3 (legacy)". Below the tabs, the input text is "The senator criticized the filibuster in Sacramento." Below the text, there are two columns: "Tokens" and "Characters". The "Tokens" column shows the value "10", and the "Characters" column shows the value "52". At the bottom, the input text is displayed with colored blocks representing individual tokens: "The" (purple), "senator" (green), "criticized" (yellow), "the" (pink), "filibuster" (blue), "in" (green), and "Sacramento." (pink).

Tokens	Characters
10	52

The senator criticized the filibuster in Sacramento.

Example: 10 tokens. “Filibuster” breaks into three pieces.


Activity, part 2: word embeddings

Same pairs. The page takes a few seconds to load its map.

 projector.tensorflow.org

1. Type “immigration” in the search box on the right. Read its nearest words.
2. Search two political words of your own.
3. Find a neighbor that is not part of the word’s meaning, like “slavery” for “immigration.” Why would Wikipedia, the map’s training text, put it there?

 **Share:** your word and that neighbor. Then: if the map had learned from censored text, what would change?



Nearest points in the original space:	
emigration	0.455
immigrants	0.528
migration	0.554
refugees	0.578
nationality	0.635
slavery	0.658
customs	0.664
taxation	0.677
immigrant	0.677

The words nearest “immigration.”

A smaller number means closer.


Screenshot: Google’s Embedding Projector, default data (Word2Vec 10K), October 6, 2026.

Training text and politics

What Chinese censors delete

- ▶ Researchers copied posts from nearly 1,400 Chinese sites before censors acted.
- ▶ Later they checked which posts were deleted, and compared the two piles across 85 topics.
- ▶ Criticism of the state, its leaders, or its policies: **no more likely** to be deleted than other posts.
- ▶ Posts about events that could bring people out together (protests, gatherings): **deleted far more often**, whatever they said about the state (about 80% on average in three such events).

Figure 9. Content of All Censored Posts (Regardless of Topic Area)




Share of posts deleted, all events. Criticism and support of the state are deleted at similar rates (no significant difference).

Source: King, Pan, and Roberts (2013), *American Political Science Review* 107(2), pp. 326 and 337; Figures 8 and 9, p. 337.

Word embeddings from a censored encyclopedia


- ▶ Two training texts: **Baidu Baike** (censored, inside China) and Chinese Wikipedia (blocked in China). One word map learned from each.
- ▶ Each row is a group of words, such as words about democracy.
- ▶ On each map: does the group sit closer to “stability” words or to “chaos” words?
- ▶ Each bar: Baidu Baike’s answer minus Wikipedia’s.
 - Red: tilted more toward chaos than on Wikipedia’s map.
 - Blue: tilted more toward stability.



Yang and Roberts (2021), Table 2. Bar length: a standardized effect size. Lighter bars: weaker evidence of a difference ($p > 0.05$).

Chatbots as voting advice: the Netherlands, 2025


- ▶ Before the October 2025 election, the Dutch data-protection authority tested four chatbots: ChatGPT, Gemini, Mistral, and Grok.
- ▶ It made up 3,000 voters, 200 for each of the 15 parties in parliament. Each voter's answers to 30 statements copied one party's, with a few changed at random.
- ▶ It asked for the three parties that best fit each voter.
- ▶ With equal voters per party, each party should come first about 6.7% of the time.
- ▶ Two parties got 55% of the first places. Several got almost none.




Source: Autoriteit Persoonsgegevens (2025), "AI-chatbots als stembulp," pp. 2 and 4 (Box 1), Figure 1, redrawn; eight smaller bars traced from the figure.

Where the chatbots sent the voters

- ▶ Left: where a Dutch voting-advice website places the parties. Several sit near the center.
- ▶ Right: where the chatbots' advice landed. Most voters on the left were sent to GroenLinks-PvdA, most on the right to the PVV. The center almost vanished.
- ▶ Every chatbot showed this pattern. The authority: do not use chatbots for voting advice.



The parties (Kieskompas)



The chatbots' advice

Axes: Links = left, Rechts = right, Progressief = progressive, Conservatief = conservative.


Source: Autoriteit Persoonsgegevens (2025), "AI-chatbots als stembulp," pp. 3–4, Figures 3 and 4 (Figure 3 from Kieskompas).

Discuss



1. What would an AI trained mostly on one country's text treat as normal? Give one example.
2. Which languages and communities are hard for AI to read, and who pays for that?
3. Back to today's question: if a model learned from censored text, how could you tell?

Next week



Next Wednesday, October 14

Large language models.

Backup slides

A map of meaning




Illustration: positions placed by hand.

Words that fit the same blanks

Sentence with a blank

Words that fit

The bill was vetoed by the _____.

governor, president, mayor

Early this morning she cast her _____.

ballot, vote

The election was won by the _____.

party, candidate

Everyone danced at the _____.

party, wedding


On the map: king – man + woman $\approx$ queen

Source: Mikolov, Chen, Corrado, and Dean (2013); Mikolov, Yih, and Zweig (2013); the 2025 course slides.

One word, many meanings

Word	One sentence	Another sentence
party	She voted for the other party.	We danced all night at the party.
bank	The bank raised its interest rates.	They sat on the bank of the river.
right	Voting is a right.	The right won the Senate.


Attention



- ▶ Line width: how much “party” draws on each earlier word.
- ▶ See real attention in your browser: poloclub.github.io/transformer-explainer

Illustration: line widths drawn by hand. In chatbot models, a word draws only on the words before it.

From rules to large language models



Source: the 2025 course slides; Weizenbaum (1966); Jurafsky and Martin, ch. 3; Mikolov, Chen, et al. (2013); Vaswani et al. (2017); Devlin et al. (2019); Brown et al. (2020).

Sources

Autoriteit Persoonsgegevens (2025). "AI-chatbots als stembulp." *Rapportage AI & Algoritmes Nederland*, special edition, October 21.

Brown, T. B., et al. (2020). "Language Models are Few-Shot Learners." *NeurIPS* 33.

Caliskan, A., Bryson, J. J., and Narayanan, A. (2017). "Semantics derived automatically from language corpora contain human-like biases." *Science* 356(6334): 183–186.

Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." *NAACL-HLT*: 4171–4186 (arXiv 2018).

Grimmer, J., and Stewart, B. M. (2013). "Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts." *Political Analysis* 21(3): 267–297.

King, G., Pan, J., and Roberts, M. E. (2013). "How Censorship in China Allows Government Criticism but Silences Collective Expression." *American Political Science Review* 107(2): 326–343.

Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013). "Efficient Estimation of Word Representations in Vector Space." arXiv:1301.3781.

Mikolov, T., Yih, W., and Zweig, G. (2013). "Linguistic Regularities in Continuous Space Word Representations." *NAACL-HLT*: 746–751.

Petrov, A., La Malfa, E., Torr, P. H. S., and Bibi, A. (2023). "Language Model Tokenizers Introduce Unfairness Between Languages." *NeurIPS* 36.

Smilkov, D., et al. (2016). "Embedding Projector: Interactive Visualization and Interpretation of Embeddings." arXiv:1611.05469.

United Nations (1948). *Universal Declaration of Human Rights*. English and Chinese: UN texts; others: UDHR in Unicode.

Vaswani, A., et al. (2017). "Attention Is All You Need." *NeurIPS* 30: 5998–6008.

Weizenbaum, J. (1966). "ELIZA: A computer program for the study of natural language communication between man and machine." *Communications of the ACM* 9(1): 36–45.

Yang, E., and Roberts, M. E. (2021). "Censorship of Online Encyclopedias: Implications for NLP Models." *Proceedings of ACM FAccT '21*: 537–548.

Prep: Financial Times (2023), "Generative AI exists because of the transformer"; 3Blue1Brown (2024), "Transformers, the tech behind LLMs".

Icons: Font Awesome Free 5 (CC BY 4.0). Video: Manim Community Edition; narration by Kokoro-82M, a synthetic voice; tokens from OpenAI's Whisper tokenizer, run locally.